


```

In [ ]: import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University, and modified
for ece239as at UCLA.
"""

class KNN(object):

    def __init__(self):
        pass

    def train(self, X, y):
        """
        Inputs:
        - X is a numpy array of size (num_examples, D)
        - y is a numpy array of size (num_examples, )
        """
        self.X_train = X
        self.y_train = y

    def compute_distances(self, X, norm=None):
        """
        Compute the distance between each test point in X and each training point
        in self.X_train.

        Inputs:
        - X: A numpy array of shape (num_test, D) containing test data.
        - norm: the function with which the norm is taken.

        Returns:
        - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
            is the Euclidean distance between the ith test point and the jth training
            point.
        """
        if norm is None:
            norm = lambda x: np.sqrt(np.sum(x**2))
            #norm = 2

        num_test = X.shape[0]
        num_train = self.X_train.shape[0]
        dists = np.zeros((num_test, num_train))
        for i in np.arange(num_test):

            for j in np.arange(num_train):
                # ===== #
                # YOUR CODE HERE:
                # Compute the distance between the ith test point and the jth
                # training point using norm(), and store the result in dists[i, j].

                # ===== #

                dists[i, j] = norm(X[i] - self.X_train[j])

                # ===== #
                # END YOUR CODE HERE
                # ===== #

        return dists

```

```

def compute_L2_distances_vectorized(self, X):
    """
    Compute the distance between each test point in X and each training point
    in self.X_train WITHOUT using any for loops.

    Inputs:
    - X: A numpy array of shape (num_test, D) containing test data.

    Returns:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      is the Euclidean distance between the ith test point and the jth training
      point.
    """
    num_test = X.shape[0]
    num_train = self.X_train.shape[0]
    dists = np.zeros((num_test, num_train))

    # ===== #
    # YOUR CODE HERE:
    # Compute the L2 distance between the ith test point and the jth
    # training point and store the result in dists[i, j]. You may
    # NOT use a for loop (or list comprehension). You may only use
    # numpy operations.
    #
    # HINT: use broadcasting. If you have a shape (N,1) array and
    # a shape (M,) array, adding them together produces a shape (N, M)
    # array.
    # ===== #

    # Basically, we do  $(a_n^T a_n - 2a_n^T b_m + b_m^T b_m)^{0.5}$  for every n and m
    # this is equal to  $(a_n^2 - 2a_n^T b_m + b_m^2)^{0.5}$ 
    # which is equal to  $(A^2 \cdot \text{shape}(n,m) + B^2 \cdot \text{shape}(n,m) - 2 * A * B^T)^{0.5}$ 
    # the shaping can be done using broadcasting
    dists = np.sqrt(np.sum(X**2, axis=1).reshape(X.shape[0], 1) + np.sum(self.X_train**2, axis=1) - 2 * np.matmul(X, self.X_train.transpose()))

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return dists

def predict_labels(self, dists, k=1):
    """
    Given a matrix of distances between test points and training points,
    predict a label for each test point.

    Inputs:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      gives the distance between the ith test point and the jth training point.

    Returns:
    - y: A numpy array of shape (num_test,) containing predicted labels for the
      test data, where y[i] is the predicted label for the test point X[i].
    """
    num_test = dists.shape[0]
    y_pred = np.zeros(num_test)
    for i in np.arange(num_test):
        # A list of length k storing the labels of the k nearest neighbors to
        # the ith test point.
        closest_y = []

```

```

# ===== #
# YOUR CODE HERE:
#     Use the distances to calculate and then store the labels of
#     the k-nearest neighbors to the ith test point. The function
#     numpy.argsort may be useful.
#
#     After doing this, find the most common label of the k-nearest
#     neighbors. Store the predicted label of the ith training example
#     as y_pred[i]. Break ties by choosing the smaller label.
# ===== #

    top_k_closest = np.array([self.y_train[index] for index in np.argsort(dists
[i])[:k]]) # classes of top k closest
    labels, counts = np.unique(top_k_closest, return_counts=True)
    label_to_count = dict(zip(labels, counts))
    y_pred[i] = max(label_to_count.keys(), key=(lambda key: label_to_count[key
]))

# ===== #
# END YOUR CODE HERE
# ===== #

return y_pred

```