


```

In [ ]: import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University, and modified
for ece239as at UCLA.
"""
class SVM(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the SVM. Note that it has shape (C, D)
        where C is the number of classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims)

    def loss(self, X, y):
        """
        Calculates the SVM loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # compute the loss and the gradient
        num_classes = self.W.shape[0]
        num_train = X.shape[0]
        loss = 0.0

        for i in np.arange(num_train):
            # ===== #
            # YOUR CODE HERE:
            # Calculate the normalized SVM loss, and store it as 'loss'.
            # (That is, calculate the sum of the losses of all the training
            # set margins, and then normalize the loss by the number of
            # training examples.)
            # ===== #

            loss += np.sum(np.maximum([0], 1 + np.matmul(self.W, X[i]) - np.matmul(self
.W[y[i]], X[i]))) - 1
            loss /= num_train
            # ===== #
            # END YOUR CODE HERE
            # ===== #

        return loss

    def loss_and_grad(self, X, y):
        """
        Same as self.loss(X, y), except that it also returns the gradient.

```

```

        Output: grad -- a matrix of the same dimensions as W containing
               the gradient of the loss with respect to W.
        """

    # compute the loss and the gradient
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    loss = 0.0
    grad = np.zeros_like(self.W)

    for i in np.arange(num_train):
        # ===== #
        # YOUR CODE HERE:
        # Calculate the SVM loss and the gradient. Store the gradient in
        # the variable grad.
        # ===== #
        loss += np.sum(np.maximum([0], 1 + np.matmul(self.W, X[i]) - np.matmul(self
        .W[y[i]], X[i]))) - 1

        indicators = np.sign(np.maximum([0], 1 + np.matmul(self.W, X[i]) - np.matmul
        l(self.W[y[i]], X[i])))
        indicators[y[i]] = 0
        grad += indicators.reshape(indicators.shape[0], 1) * X[i]
        grad[y[i]] += -np.sum(indicators.reshape(indicators.shape[0], 1) * X[i], ax
        is=0)

        # ===== #
        # END YOUR CODE HERE
        # ===== #

    loss /= num_train
    grad /= num_train

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X,y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + ab
        s(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, g
        rad_analytic, rel_error))

    def fast_loss_and_grad(self, X, y):
        """
        A vectorized implementation of loss_and_grad. It shares the same

```

```

        inputs and outputs as loss_and_grad.
        """
        loss = 0.0
        grad = np.zeros(self.W.shape) # initialize the gradient as zero

        # ===== #
        # YOUR CODE HERE:
        # Calculate the SVM loss WITHOUT any for loops.
        # ===== #
        loss = (np.sum(np.maximum([0], 1 + np.matmul(X, self.W.transpose())) - np.sum(
np.multiply(self.W[y], X), axis=1).reshape(X.shape[0], 1))) - X.shape[0]) / X.sha
pe[0]
        # ===== #
        # END YOUR CODE HERE
        # ===== #

        # ===== #
        # YOUR CODE HERE:
        # Calculate the SVM grad WITHOUT any for loops.
        # ===== #
        indicators = np.sign(np.maximum([0], 1 + np.matmul(X, self.W.transpose())) - n
p.sum(np.multiply(self.W[y], X), axis=1).reshape(X.shape[0], 1)))
        indicators[range(X.shape[0]), y] = 0
        indicators[range(X.shape[0]), y] = -np.sum(indicators, axis=1)
        grad = np.matmul(indicators.transpose(), X)/X.shape[0]
        # ===== #
        # END YOUR CODE HERE
        # ===== #

        return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
        batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
      training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
      means that X[i] has label 0 ≤ c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number
of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights
of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):

```

```

X_batch = None
y_batch = None

# ===== #
# YOUR CODE HERE:
#   Sample batch_size elements from the training data for use in
#   gradient descent. After sampling,
#   - X_batch should have shape: (dim, batch_size)
#   - y_batch should have shape: (batch_size,)
#   The indices should be randomly generated to reduce correlations
#   in the dataset. Use np.random.choice. It's okay to sample with
#   replacement.
# ===== #
indices = np.random.choice(X.shape[0], batch_size, replace=True)
X_batch = X[indices]
y_batch = y[indices]
# ===== #
# END YOUR CODE HERE
# ===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
#   Update the parameters, self.W, with a gradient step
# ===== #
self.W -= learning_rate * grad
# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[0])

    # ===== #
    # YOUR CODE HERE:
    #   Predict the labels given the training data with the parameter self.W.
    # ===== #
    y_pred = np.argsort(np.matmul(X, self.W.transpose()), axis=1)[:X.shape[0], self.W.shape[0] - 1]
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```

